

Chi-Squared Test (Test of Independence)

A chi-squared test (χ^2 test) is used to decide whether there is a relationship exists between two categorical variables of a sample.

In this exercise, we will perform χ^2 test based on 'Immunotherapy.csv' and 'titanic.csv' dataset.

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from scipy import stats
```

Immunotherapy Dataset

```
In [2]: df_immuno = pd.read_csv("https://raw.githubusercontent.com/ThammakornS/ProgStat/main/Immunotherapy.csv")
df_immuno.head()
```

```
Out[2]:   gender  age  time  number_of_warts      type  area  induration_diameter  response
0  Female    15   1.75                  1  Plantar    49                 7        No
1  Female    38   2.50                  1     Both    43                50       Yes
2  Female    24   4.25                  1  common   174                30       Yes
3  Female    34   8.50                  1  Plantar   163                 7        No
4  Female    53  10.00                  1  Plantar    30                25       Yes
```

Adjust Data Type

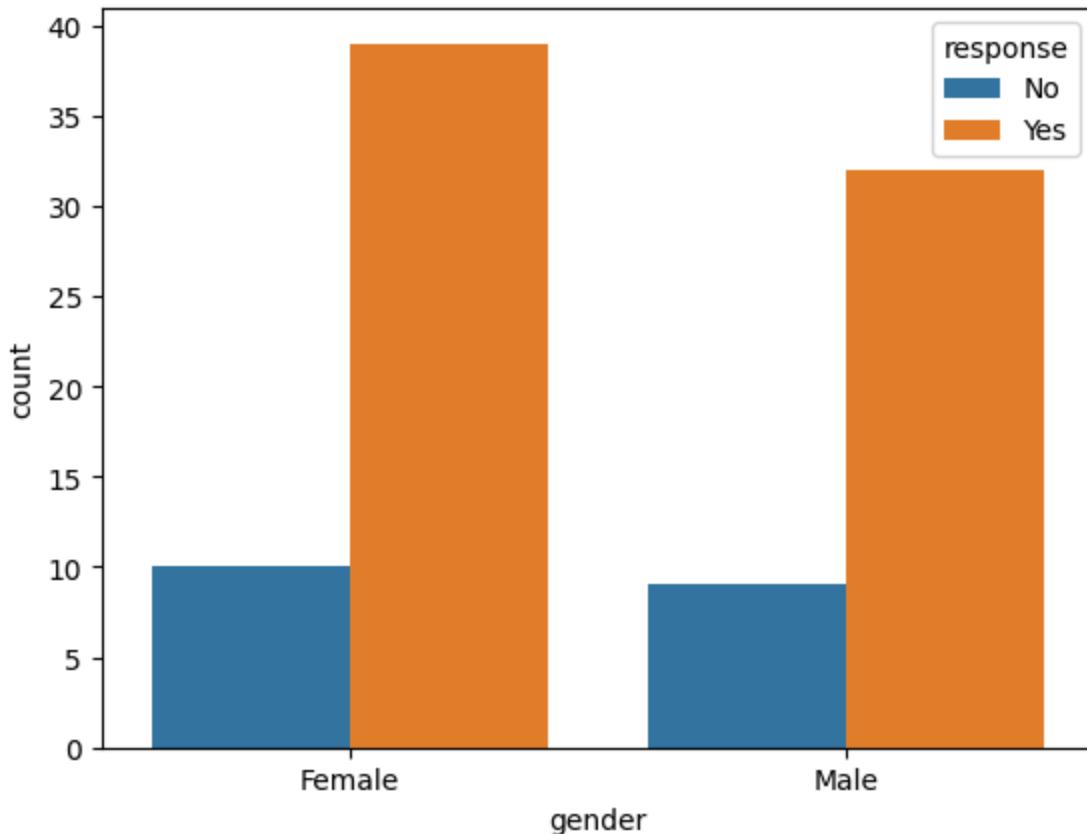
```
In [3]: df_immuno = df_immuno.astype({
    'gender': 'category',
    'type': 'category',
    'response': 'category'
})
df_immuno.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 90 entries, 0 to 89
Data columns (total 8 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   gender          90 non-null     category
 1   age              90 non-null     int64   
 2   time             90 non-null     float64
 3   number_of_warts  90 non-null     int64   
 4   type             90 non-null     category
 5   area             90 non-null     int64   
 6   induration_diameter  90 non-null  int64   
 7   response         90 non-null     category
dtypes: category(3), float64(1), int64(4)
memory usage: 4.3 KB
```

Visualization of Categorical Data

```
In [4]: sns.countplot(data=df_immuno,
                     x='gender',
                     hue='response')
```

```
Out[4]: <Axes: xlabel='gender', ylabel='count'>
```



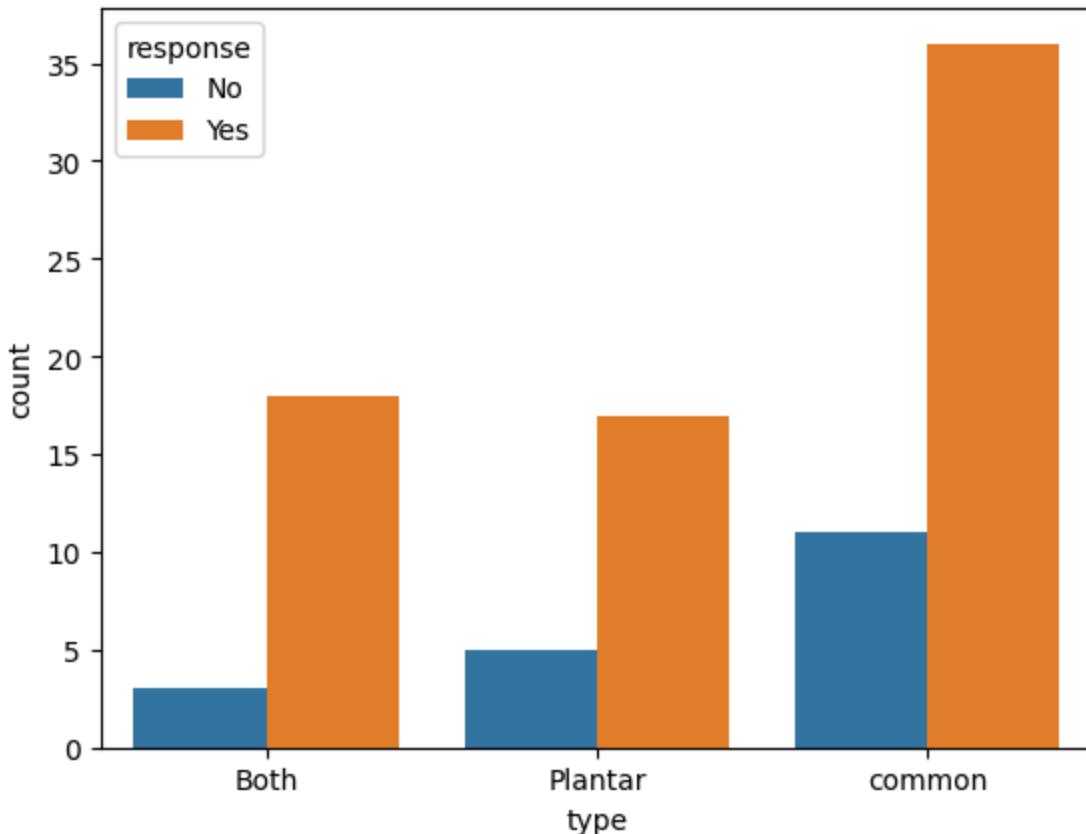
```
In [5]: pd.crosstab(df_immuno.gender,
                    df_immuno.response,
                    margins=True)
```

```
Out[5]: response No Yes All
```

gender	No	Yes	All
Female	10	39	49
Male	9	32	41
All	19	71	90

```
In [6]: sns.countplot(data=df_immuno,  
                     x='type',  
                     hue='response')
```

```
Out[6]: <Axes: xlabel='type', ylabel='count'>
```



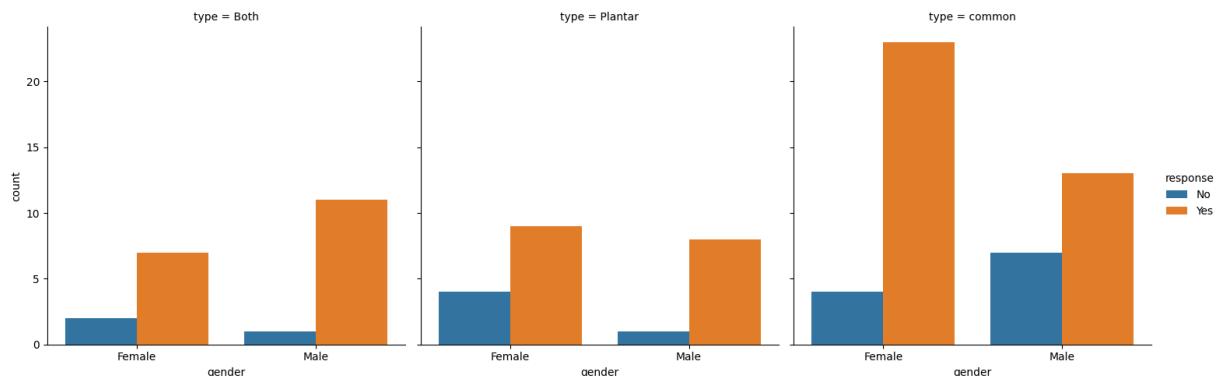
```
In [7]: pd.crosstab(df_immuno.type,  
                     df_immuno.response,  
                     margins=True)
```

```
Out[7]: response No Yes All
```

type	3	18	21
Both	3	18	21
Plantar	5	17	22
common	11	36	47
All	19	71	90

```
In [8]: sns.catplot(data=df_immuno,  
                  kind='count',  
                  col='type',  
                  x='gender',  
                  hue='response')
```

```
Out[8]: <seaborn.axisgrid.FacetGrid at 0x20bfc4c5400>
```



Gender & Response

Create Contingency Table

To run the Chi-Square Test, the easiest way is to convert the data into a contingency table with frequencies.

```
In [9]: contab = pd.crosstab(df_immuno.gender,  
                           df_immuno.response)  
contab
```

```
Out[9]: response No Yes
```

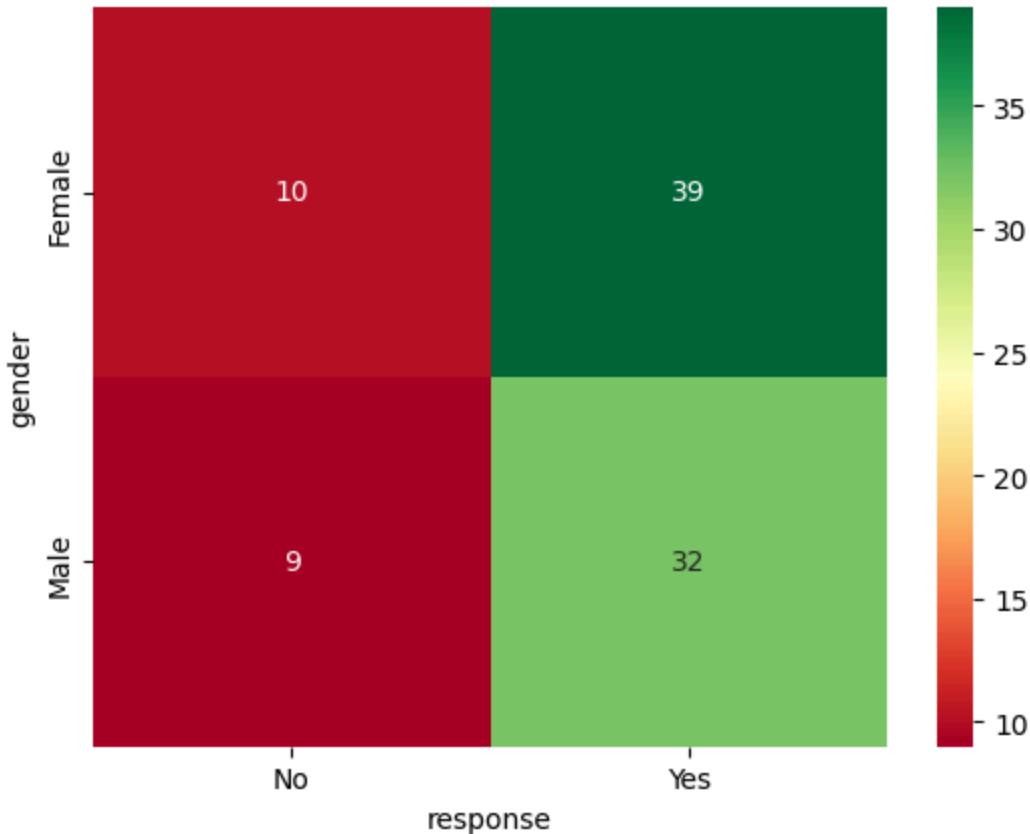
gender	10	39
Female	10	39
Male	9	32

Visualize Contingency Table

An easy way to visualize a contingency table is to draw a heatmap.

```
In [10]: sns.heatmap(contab, annot=True, cmap='RdYlGn')
```

```
Out[10]: <Axes: xlabel='response', ylabel='gender'>
```



Chi-Square Test

Now that we have built the contingency table, we can pass it to `chi2_contingency()` from the `scipy` library which returns:

- The test statistic (`c`)
- The p-value of the test (`p`)
- Degrees of freedom (`dof`)
- The expected frequencies, based on the marginal sums of the table (`expected`)

Hypothesis

H_0 : the gender & response have *no relationship*

H_A : there is a relationship between gender & response

```
In [11]: c, p, dof, expected = stats.chi2_contingency(contab)
```

```
In [12]: print("p-value:", p)
```

p-value: 1.0

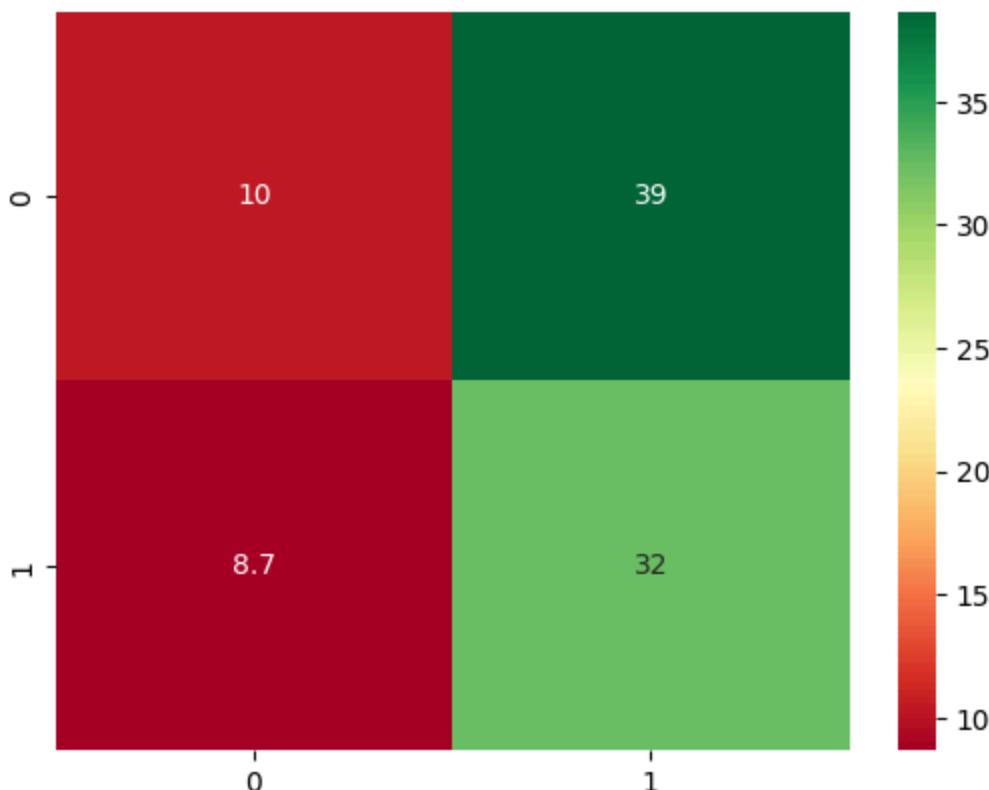
p-value > 0.05, fail to reject H_0 .

The p-value is 1 which means that we do not reject the H_0 at 95% level of confidence.

Expected frequencies (no relationship):

```
In [13]: sns.heatmap(expected, annot=True, cmap='RdYlGn')
```

```
Out[13]: <Axes: >
```



The expected value at row i and column j can be calculated using below equation.

$$\frac{R_i \text{ Total} \times C_j \text{ Total}}{\text{Grand Total}}$$

Let's see the contingency table again before calculating expected values.

```
In [14]: pd.crosstab(df_immuno.gender,
                     df_immuno.response,
                     margins=True)
```

```
Out[14]: response  No  Yes  All
```

		gender		
		Female	39	49
		Male	32	41
		All	19	71
				90

Now calculate expected value at gender=Female, response=No

```
In [15]: (49*19)/90
```

```
Out[15]: 10.344444444444445
```

Calculate expected value at gender=male, response=No

```
In [16]: (41*19)/90
```

```
Out[16]: 8.655555555555555
```

Type & Response

Create Contingency Table

```
In [17]: contab = pd.crosstab(df_immuno.type,
                             df_immuno.response)
contab
```

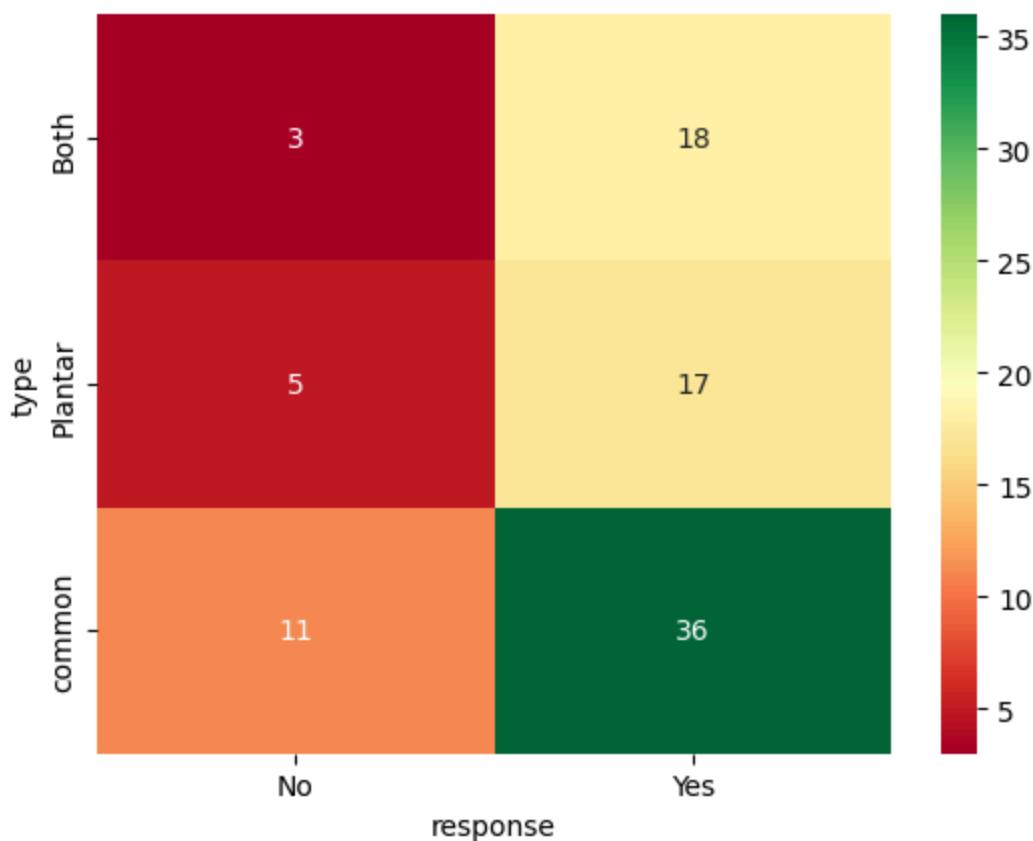
```
Out[17]: response  No  Yes
```

		type	
		Both	18
		Plantar	17
		common	36
			11

Visualize Contingency Table

```
In [18]: sns.heatmap(contab, annot=True, cmap='RdYlGn')
```

```
Out[18]: <Axes: xlabel='response', ylabel='type'>
```



Chi-Square Test

Hypothesis

H_0 : the type & response variables have *no relationship*

H_A : there is a relationship between type & response variables

```
In [19]: c, p, dof, expected = stats.chi2_contingency(contab)
```

```
In [20]: print("p-value:", p)
```

p-value: 0.6803408056744953

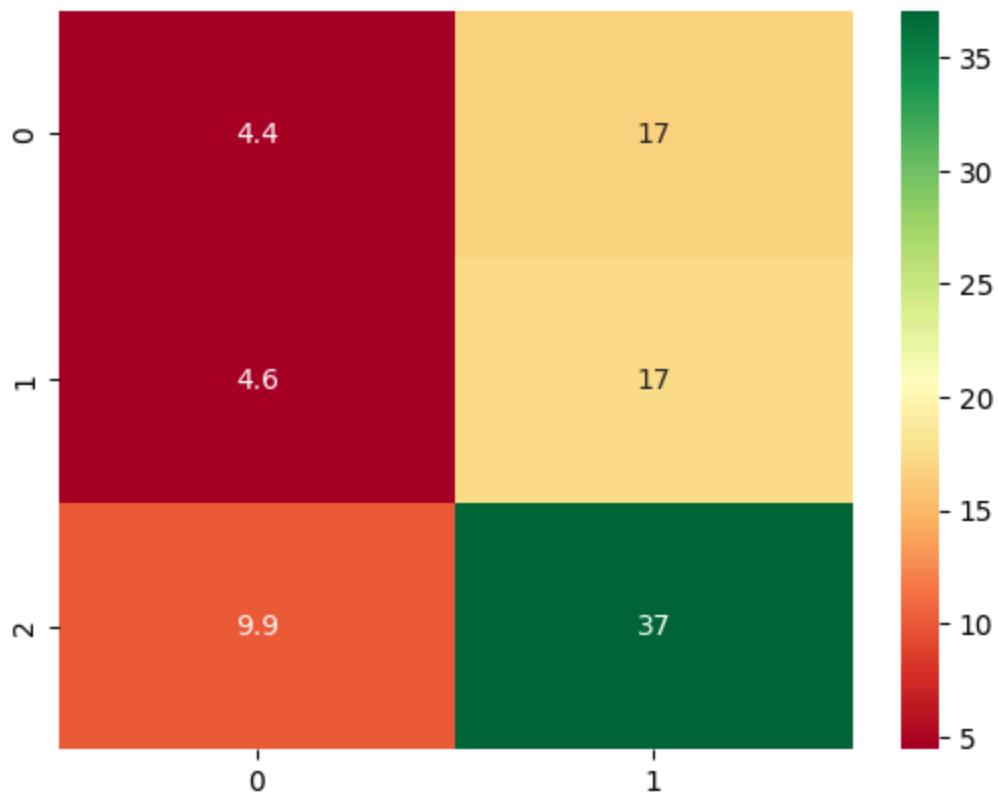
p-value > 0.05, fail to reject H_0 .

The p-value is 1 which means that we do not reject the H_0 at 95% level of confidence.

Expected frequencies:

```
In [21]: sns.heatmap(expected, annot=True, cmap='RdYlGn')
```

```
Out[21]: <Axes: >
```



Titanic Dataset

```
In [22]: df_titan = pd.read_csv("https://raw.githubusercontent.com/ThammakornS/ProgStat/main/titanic.csv")
df_titan.head()
```

Out[22]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...)	female	38.0	1	0	PC 17599	71.2833
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500

Adjust Data Type

In [23]:

```
df_titan = df_titan.astype({
    'Survived': 'category',
    'Pclass': 'category',
    'Sex': 'category',
    'Cabin': 'category',
    'Embarked': 'category'
})
df_titan.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   PassengerId 891 non-null    int64  
 1   Survived     891 non-null    category
 2   Pclass       891 non-null    category
 3   Name         891 non-null    object  
 4   Sex          891 non-null    category
 5   Age          714 non-null    float64 
 6   SibSp        891 non-null    int64  
 7   Parch        891 non-null    int64  
 8   Ticket       891 non-null    object  
 9   Fare          891 non-null    float64 
 10  Cabin         204 non-null    category
 11  Embarked     889 non-null    category
dtypes: category(5), float64(2), int64(3), object(2)
memory usage: 59.8+ KB
```

Gender & Survived

Create Contingency Table

```
In [24]: contab = pd.crosstab(df_titan.Sex,
                               df_titan.Survived,)

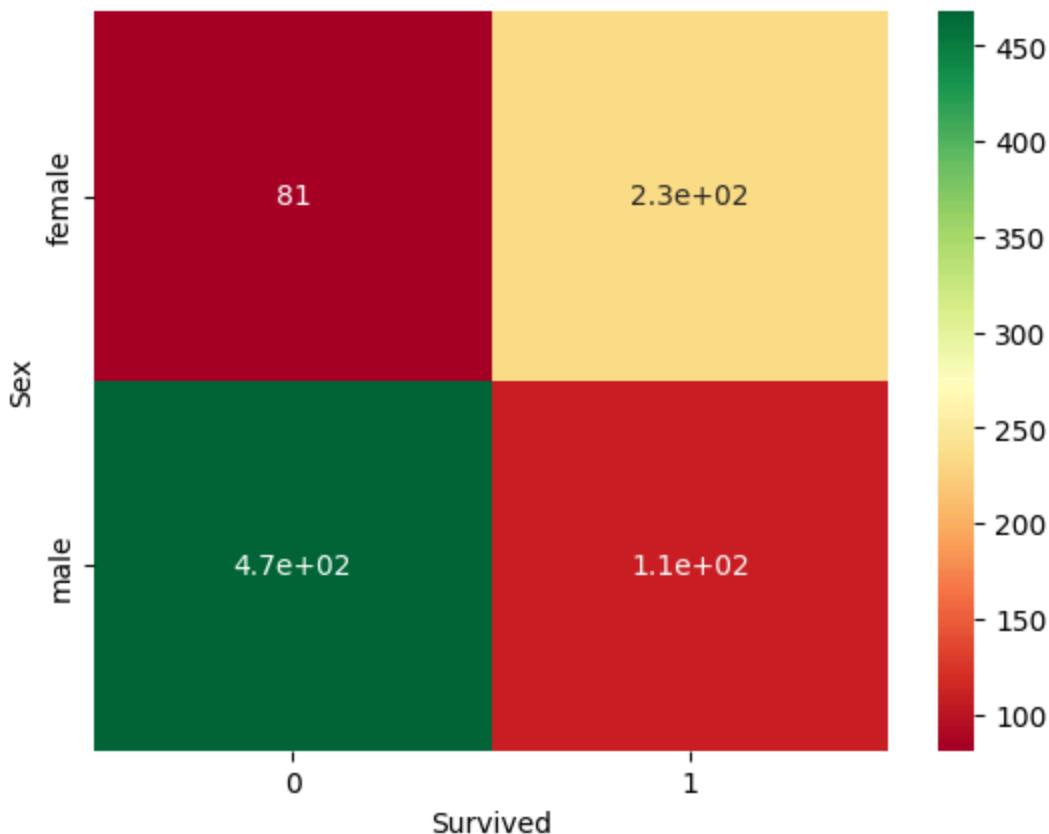
contab
```

```
Out[24]: Survived    0    1
          Sex
          female  81  233
          male   468 109
```

Visualize Contingency Table

```
In [25]: sns.heatmap(contab, annot=True, cmap='RdYlGn')
```

```
Out[25]: <Axes: xlabel='Survived', ylabel='Sex'>
```



Chi-Square Test

Hypothesis

H_0 : the gender & survived variables have *no relationship*

H_A : there is a relationship between gender & survived variables

```
In [26]: c, p, dof, expected = stats.chi2_contingency(contab)
```

```
In [27]: print("p-value:", p)
```

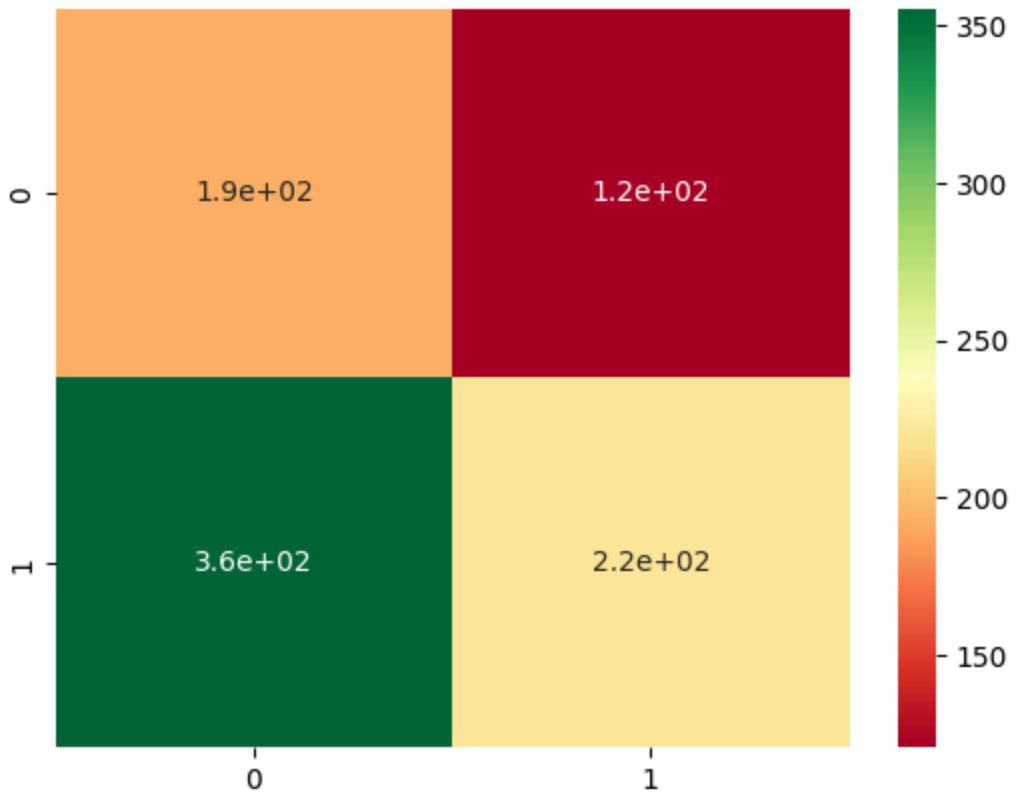
p-value: 1.1973570627755645e-58

p-value <= 0.05, reject H_0 , accept H_A .

Expected frequencies:

```
In [28]: sns.heatmap(expected, annot=True, cmap='RdYlGn')
```

```
Out[28]: <Axes: >
```



Pclass & Survived

Create Contingency Table

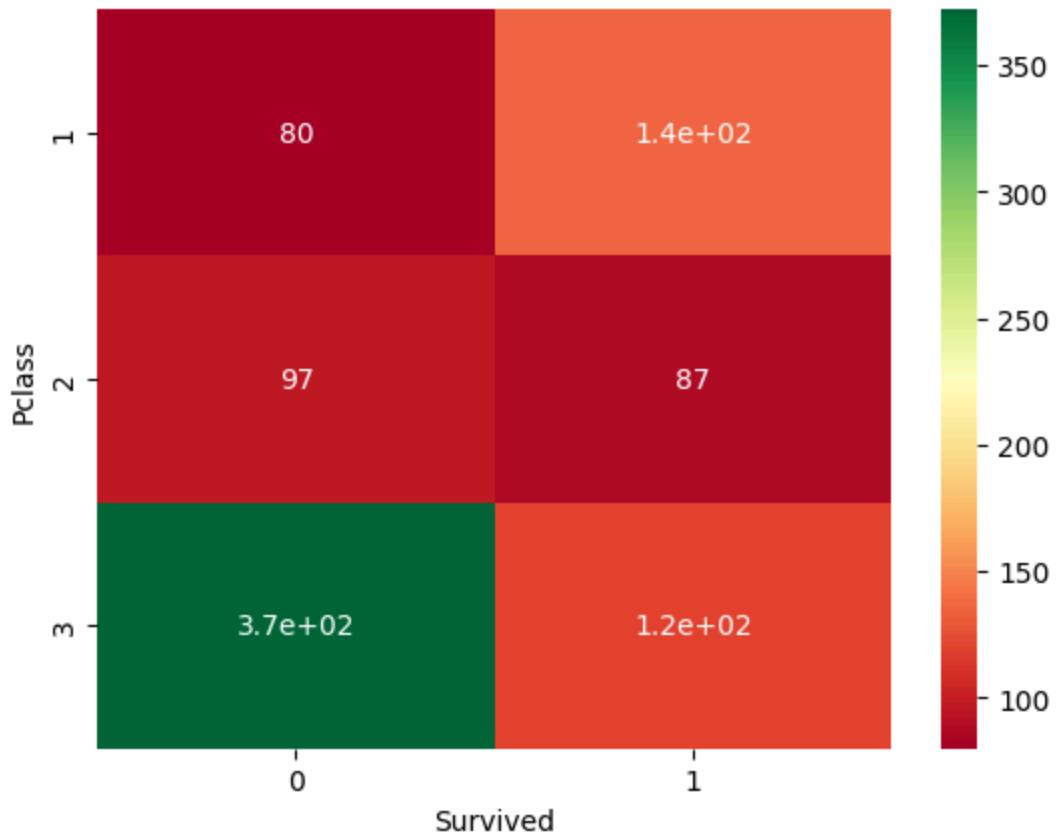
```
In [29]: contab = pd.crosstab(df_titan.Pclass,
                             df_titan.Survived)
contab
```

```
Out[29]: Survived    0    1
Pclass
   1    80  136
   2    97   87
   3   372  119
```

Visualize Contingency Table

```
In [30]: sns.heatmap(contab, annot=True, cmap='RdYlGn')
```

```
Out[30]: <Axes: xlabel='Survived', ylabel='Pclass'>
```



Chi-Square Test

Hypothesis

H_0 : the pclass & survived variables have *no relationship*

H_A : there is a relationship between pclass & survived variables

```
In [31]: c, p, dof, expected = stats.chi2_contingency(contab)
```

```
In [32]: print("p-value:", p)
```

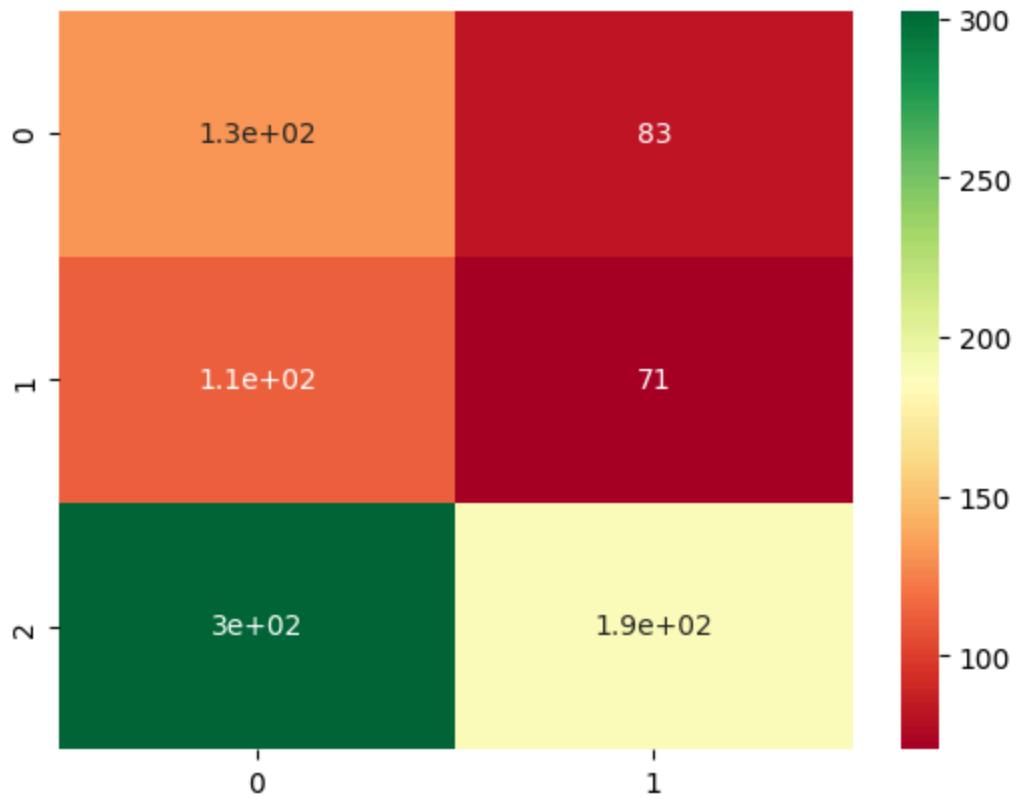
p-value: 4.549251711298793e-23

p-value <= 0.05, reject H_0 , accept H_A .

Expected frequencies:

```
In [33]: sns.heatmap(expected, annot=True, cmap='RdYlGn')
```

```
Out[33]: <Axes: >
```



In []: